

AI Safety Interview Scoring Rubric

Based on roles listed on the 80,000 Hours AI Safety job board, June 2026

Role Type	What they look for	# Criteria	Typical orgs
Technical Research	ML engineering, interpretability, alignment, model evals	10	Anthropic, OpenAI, CAIS, FAR AI, DeepMind
Security & Red-Teaming	Adversarial testing, threat modeling, dangerous capability evals	9	Anthropic, OpenAI Preparedness, UK AISI
Policy	Governance analysis, regulatory engagement, technical literacy	9	GovAI, CAIS, Anthropic Policy, UK AISI
Operations	Programme management, hiring, org design, CoS support	9	CAIS, GovAI, BlueDot Impact
Research Economics	Econometrics, labour economics, AI impacts research	9	Anthropic (London/Singapore)
Field-Building & Education	Curriculum design, fellowships, community building, mentorship	9	BlueDot Impact, Apart Research, CAIS
Cultural Fit (all roles)	Adversarial grit, epistemic rigour, strategic foresight, conceptual engineering	4	Applies to every role · scored separately

Scoring Scale
1 – No exposure
2 – Aware but untested
3 – Some experience
4 – Solid foundation
5 – Demonstrated strength

How to use: This is an interviewer scoring tool. For each candidate, open the relevant role tab AND the Cultural Fit tab. Score the candidate 1–5 on each criterion using the Score 1 / 3 / 5 descriptors as anchors, and capture evidence in the notes column. Each criterion has a Weight (3 = core, 2 = important, 1 = useful) that you can edit per role or per hire. The tab auto-calculates a Weighted Score (1–5) and a Readiness % from only the criteria you have scored. Cultural Fit produces its own separate score and is deliberately not blended into the role score: treat skill as the floor and the traits as the ceiling, and read the two numbers together. Low scores (1–2) flag red and strengths (4–5) flag green automatically.

Weight key

3 – Core: a primary screen for this role

2 – Important: meaningful but not disqualifying

1 – Useful: a plus, not a gate